From Curiosity to Creation: Build Your First AI Project & Research Portfolio

Prof. Ping Chen   Inal Mashukov

Department of Computer Science
University of Massachusetts Boston

DAY 4 - Vibes, Aesthetics & Creative AI

Summer 2026

Discriminative vs. Generative AI

Discriminative Models (The Past)

These models learn the boundary *between* classes. They are used for classification and prediction.

- **Input:** An image of a dog.
- **Output:** The label "Dog".
- **Goal:** $P(Y|X)$ - Predict the label given the data.

Generative Models (The Present)

These models learn the underlying distribution of the data itself. They create entirely new data.

- **Input:** The word "Dog".
- **Output:** A brand new, photorealistic image of a dog that has never existed.
- **Goal:** $P(X|Y)$ - Generate the data given the label.

Can a Machine Be Creative?

The Human View of Creativity

- We view art, music, and storytelling as deeply emotional, subjective human experiences.
- Is an AI creating art, or just committing high-tech plagiarism by blending its training data?

The Algorithmic View

"*Creativity is just connecting things.*" - Steve Jobs.

If creativity is simply taking two existing concepts and interpolating a novel bridge between them, neural networks are mathematically perfectly suited for it.

What is the Latent Space?

The Vibe Landscape

Generative models compress all human visual concepts into a massively high-dimensional geometric map called the **Latent Space**.

- It is a coordinate system of concepts.
- Every possible image is a single $(x, y, z, \ldots)$ coordinate in this space.
- If you move slightly to the left, the image might get "happier." If you move up, it might get "more photorealistic."

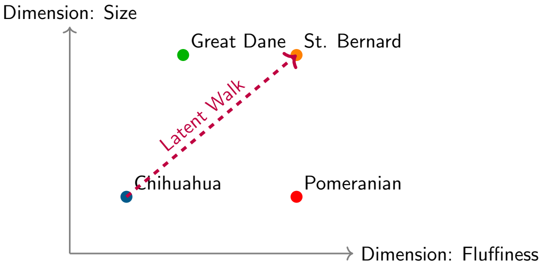

Latent Interpolation: Math Becomes Art

Vector Algebra with Images

Because the Latent Space is continuous, we can pick any two points and mathematically walk between them.

- Point A: A photograph of an apple.
- Point B: A photograph of a dog.
- The exact midpoint $(A+B)/2$ is a smooth, AI-generated hybrid "apple-dog."

Prompting is just GPS navigation

When you type a prompt into Midjourney or DALL-E, you are not telling it how to draw. You are giving the algorithm a set of GPS coordinates to locate a specific point in the Latent Space, which it then renders into pixels.

How do we generate the pixels? (Diffusion Models)

Modern image AI (Stable Diffusion, DALL-E 3) relies on a physics concept called **Diffusion**.

Phase 1: Forward Diffusion (Destruction)

During training, the AI takes a clean image of a cat and slowly adds static (Gaussian noise) to it over 1,000 steps until it is completely unrecognizable TV static.

Phase 2: Reverse Diffusion (Creation)

The neural network is trained to do the exact opposite. We give it pure static, and teach it to subtract the noise, step-by-step, until a clean image emerges.

Visualizing the Diffusion Process

From Pure Noise to Structured Art

Text Conditioning Guides the Noise Removal

Text Prompt: "A cute dog"

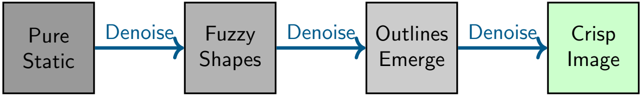

Neural Style Transfer (NST)

Separating Content from Style

AI researchers discovered that Convolutional Neural Networks (CNNs) process images in a hierarchy that allows us to mathematically separate *what* is in an image from *how* it is drawn.

- **Content:** The physical objects (e.g., a house, a tree). These are recognized in the *deep* layers of the network.
- **Style:** The brushstrokes, colors, and textures. These are captured in the *shallow* layers.

The Formula for Art

By combining the Content vector of a photograph with the Style vector of a painting (e.g., Van Gogh's *Starry Night*), the network optimizes a new image that perfectly blends both.

Textual "Vibe Transfer"

We can apply the exact same concept to language using Large Language Models. We can separate the **semantic content** (the facts) from the **syntactic style** (the vibe).

- *Base Content:* "The AI model achieved 95% accuracy on the test set."
- *Pirate Style:* "Arrr, the metal brain conquered 95% of the treacherous exam!"
- *Academic Style:* "Empirical evaluation demonstrates a 95% top-1 accuracy metric across the validation distribution."

The Mechanics of Vibe

This works because the LLM traverses the latent space. By anchoring the generation to a specific contextual cluster (e.g., 18th-century nautical terms), it forces the output probabilities to align with that aesthetic.

Multi-Agent Systems: AI Talking to AI

A single AI model has limitations. It gets stuck in loops or hallucinates without realizing it.

The Solution: Multi-Agent Frameworks

- Instead of using one LLM, we initialize two (or more) separate AI personas and force them to talk to each other.
- **Agent 1 (The Generator):** Tasked with brainstorming and creating raw ideas.
- **Agent 2 (The Critic):** Tasked with finding logical flaws, plot holes, or aesthetic errors.

Emergent Creativity

By looping their conversation, the system self-corrects. The final output is vastly superior to what a single agent could produce on its own.

Visualizing the Collaborative Loop

Iterative Feedback Loop (Max N turns)

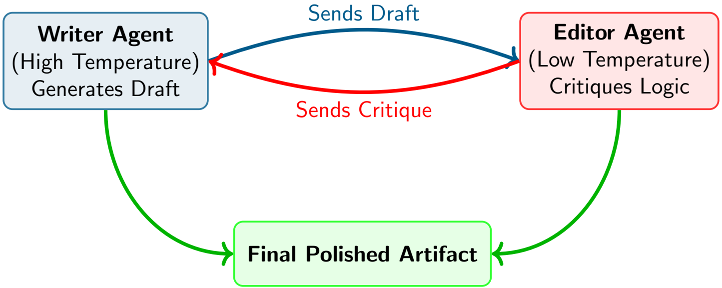

Generative AI Safety & Ethics

With the ability to perfectly synthesize human text, audio, and images, Generative AI introduces massive ethical risks.

- **Deepfakes:** Non-consensual generation of realistic photos/audio of real people.
- **Copyright:** Generative models are trained on billions of scraped images from human artists who were not compensated.
- **Misinformation:** LLMs can generate fake news articles at infinite scale, complete with fabricated citations.

The Alignment Problem

How do we program human values into a matrix of floating-point numbers? How do we ensure the model refuses dangerous prompts (e.g., "How to build a bomb") without censoring harmless creativity?

The Irony of Creativity

"Creativity is just highly-rewarded hallucination."

When we want facts:
If you ask an AI for the capital of France and it says "A shimmering city of glass," that is a **hallucination**. It is a failure of the system.

When we want art:
If you ask an AI to write a fantasy novel and it describes "A shimmering city of glass," that is **creativity**. It is a success.

Mechanistically, the AI is doing the exact same mathematical operation in both scenarios.

PhD-Led Lab 1: Vibe Transmutation (1:00 PM)

Image Style Mixing

- We will access a generative image pipeline.
- You will experiment with prompt weighting (e.g., `dog:1.5, cyberpunk:-0.5`) to mathematically navigate the latent space.
- Goal: Force two conflicting concepts to merge seamlessly.

Text "Vibe Transfer"

- We will take a dry, factual Wikipedia article.
- You will engineer prompts that force the LLM to rewrite the article in increasingly bizarre constraints (e.g., as a Shakespearean sonnet, as 90s hacker slang).

PhD-Led Lab 2: Multi-Agent Story Creation (2:40 PM)

Building the Collaboration Loop

- You will programmatically link two LLM instances using Python.
- **Agent A** will be instructed to write a short story.
- **Agent B** will be instructed to ruthlessly critique Agent A's story.
- You will run the loop 3 times and analyze how the artifact evolves.

The Output

This lab forms the foundation for advanced AI architectures. You are no longer just prompting; you are building automated software ecosystems.

Your Goal for Today (4:00 PM Check-In)

The Deliverable
You must finalize your **Core Project Prototype**.

- *Visual:* Generate the final artifacts or images for your portfolio.
- *Narrative:* Output the final text from your multi-agent system.
- *Interactive:* Ensure your code/chatbot runs without crashing.

Tomorrow is the Showcase

Tomorrow, you will compile everything (your research question, your interpretability explanation, your philosophical reflection, and today's final artifact) into a formal presentation.

Thank you!
Questions?

E-mail: inal.mashukov001@umb.edu